BigML Assignment 3: Streaming Phrase Finding

Due Mon, Feb 17 before class (1:29pm) via Autolab
Late submission: 50% credit before Wed, Feb 19, 1:29pm via Autolab

1 Important Note

This assignment is the first of two that use the phrase finding algorithm we discussed in class (see Tomokiyo and Hurst, 2003, http://dl.acm.org/citation.cfm?id=1119287). You will be expected to reuse the code you develop for this assignment for a future assignment, and **you are expected to use Java for this assignment**.

Please post clarification questions to the Piazza, and the instructors can be reached at the following email address: *10605-Instructors@cs.cmu.edu*.

2 Streaming Phrase Finding

This assignment is based on the method in *"A language model approach to keyphrase extraction"*, Tomokiyo and Hurst, 2003, http://dl.acm.org/citation.cfm?id=1119287 . The basic goal is to pick out word sequences in a corpus that are meaningful as phrases (in the paper's words, have high "phraseness"), and over-represented in the corpus (have high "informativeness"). Phraseness and informativeness are computed based on some simple frequency features, which we will denote as

C(x,y)	frequency of the phrase "x y" in the corpus of interest
B(x,y)	analogous for the background corpus
C(x), C(y)	frequency of a word x or y in the corpus of interest
B(x), B(y)	analogous for the background corpus

You also need a handful of constants, like the number of words and bigrams in each corpus, and the vocabulary sizes. For this assignment, the *corpus of interest* (foreground corpus) is those books published in the years 1990-1999. The background corpus is those books published in the years 1960-1989. We've already done the preprocessing for you (for example, we've thrown out words and phrases with non-letter characters, normalized the words to their lowercase variants, and aggregated the counts by decade). Please do NOT remove the stopwords. As before, there is a smaller test set for debugging your code.

We suggest using a simple data structure that associates a set of attribute-value pairs with a phrase. Once you have this data structure - for instance like the examples below - it's easy to compute phraseness and informativeness in a streaming setting:

Key	Value
united states	Cxy=104324,Bxy=214442,Cx=321313,Cy=424134,Bx=23141,By=444141
white house	Cxy=4343003,Bxy=43322,Cx=2344233,Cy=786677,Bx=235661,By=1056773

Consider the first row in the example above. Cxy denotes the C(x,y) count for the phrase `united states` in the foreground corpus. Cx is the count for `united` and Cy is the count for `states` in the foreground corpus. It's easy enough to tally up these counts over the corpus. The tricky part is getting all the counts for one bigram into one place.

As in assignment 2, we're going to assume the unigrams won't fit in memory; all Java commands must be followed with -Xmx128m. So, we will implement a stream and sort algorithm to request count, fulfill the requests, and aggregate the responses. Here's a general outline of the algorithm for aggregating the counts (refer also to the recent lecture slides):

1. Compute counter files mapping $x \rightarrow Bx, x \rightarrow Cx, xy \rightarrow Cxy, xy \rightarrow Bxy$

   - This is mostly done, but you will need to aggregate together the counts for the different decades in the background corpus.

2. Gather together the background and foreground counts for each unigram to create one data structure with key `x` and value `Bx=some number,Cx=some number`. Do the same for the bigrams.

3. Stream through all the phrases and for each phrase "x y", create two messages–one asking for the unigram frequencies for x, and one asking for the frequencies of y. Each message is just a pair consisting of the query word and the phrase making the request: so the phrase `united states` will generate the messages

   ```
   united,united states
   states,united states
   ```

4. To deliver the messages, sort them in with the unigram frequency file. Use a secondary key so the attribute-value pairs come first, and the messages come last (as described in lecture). While scanning through the unigram file, you will do something like this:

   (a) for each distinct key x (e.g. `united`)
      i. read the attribute-value pairs for x (e.g. `Bx=14342,Cx=24123`)

ii. for each message from step two addressed to x from a phrase of the form xz (e.g. `united states`):

    A. send a message to xz with the attribute-value list at the content. (i.e. write out a pair with key xz and value of the attribute-value list (eg, `united states, Bx=14342,Cx=24123`)

iii. for each message from step two addressed to x from a phrase of the form zx (e.g. `fly united`):

    A. send a message to zx with the attribute-value list at the content. (i.e. write out a pair with key zx and value of the attribute-value list (eg, `united states, By=14342,Cy=24123`)

    (***Notice the change from x to y in the value list***)

5. Now you've created two new data structures for each bigram, one with `Bx=some number,Cx=some number` and one with `By=some number,Cy=some number` . To deliver these messages, sort them in with the bigram frequency file that contains `Bxy=some number,Cxy=some number` , and merge the data structures.

Now you have all of the counts in the same data structure, and can compute the phraseness and informativeness scores as described in the paper. Recall the definition of point wise KL Divergence:

$$\delta_w(p||q) = p(w)log\left(\frac{p(w)}{q(w)}\right) \tag{1}$$

Phraseness is point wise KL Divergence with

$$p = p_{fg}(x \wedge y) \quad q = p_{fg}(x)p_{fg}(y). \tag{2}$$

Informativeness is point wise KL Divergence with

$$p = p_{fg}(x \wedge y) \quad q = p_{bg}(x \wedge y) \tag{3}$$

Where p_{fg} is the probability of an event under the foreground corpus (corpus of interest), and p_{bg} is the probability of an event under the background corpus. The phrase score is just the sum of phraseness and informativeness. Please use the natural logarithm for this assignment. The paper uses a more complex smoothing algorithm, but you should just use add one smoothing, as in Assignments 1 & 2.

3   The Data

We are using the google books corpus. We've done some preprocessing of the data, and created two sets of data files. This is unsupervised learning, so there is no test set.

The data has the following format:

```
<text>\t<decade>\t<count>
```

where `text` is either a bigram or a unigram, and count is the number of times that text occurred in a book in the given decade. The smaller dataset is all bigrams that contain the word *apple*, and all unigrams appearing in those bigrams. This subsample was chosen because Apple Computer became a more popular company in the 90s, and thus the word Apple took on a different meaning. You should be able to see that trend in the phrase statistics you calculate for the smaller dataset.

The data appears at `/afs/cs.cmu.edu/project/bigML/phrases/`. Files with the word *apple* in the name are the subsampled datasets.

4 Deliverables

4.1 Report

Submit your implementations via AutoLab. You should implement the algorithm by yourself instead of using any existing machine learning toolkit. You should upload your code (including all your function files) along with a **report**, which should solve the following questions:

1. The top 20 phrases (sorted by total score) with their phraseness and informativeness scores from the full data set and the apple data set. (2 points)

2. What do you notice about the phrases ranked highest in your results for the two data sets? Do they give you any insights into events or trends in the 90s? (2 points)

3. Are there any downfalls you see to using the total phrase score? For example, are there some phrases that are ranked high even though you don't think they should be? Why are they ranked so high? (2 points)

4. How could you improve upon the total score proposed by Tomokiyo and Hurst? (2 points)

5. Consider the workflow discussed on 1/27 in class:

   ```
   % java CountForNB train.dat  > eventCounts.dat
   % java CountsByWord eventCounts.dat | sort | java CollectRecords > words.dat
   % java requestWordCounts test.dat \
     | cat - words.dat | sort | java answerWordCountRequests \
     | cat - test.dat| sort | testNBUsingRequests
   ```

   and also consider a corpus with these documents in class "breakfast":

   - *Joe likes toast*

- *Jane likes toast and jam*
- *Joe burnt his toast*

and these in class "dinner":

- *Joe likes steak*
- *Mary likes steak and ale*

and finally the test document:

- *Jane ordered eggs and toast*

Part 1. Answer the questions below (5 points):

(a) What are the entries in `eventCounts.dat` associated with the words "toast", "likes", and "steak"?

(b) What are the entries in `words.dat` associated with the words "toast", "likes", and "steak"?

(c) What is the output of `requestWordCounts` on the test corpus? Please write key values pairs as "key//value" so we can see the different parts easily.

(d) What is the output of `answerWordCountRequests` on the test corpus?

(e) What is the input to `testNBUsingRequests`?

Part 2. Suppose there are K classes, V distinct words in the training corpus, and N tokens in the test corpus. Answer the questions below (7 points):

(a) The number of integers that are stored in `eventCounts.dat`.

(b) The number of key-value pairs that are stored in `eventCounts.dat`.

(c) The number of integers that are stored in `words.dat`.

(d) The number of key-values pairs that are stored in `words.dat`.

(e) The number of key-value pairs output by `requestWordCounts`.

(f) The number of key-value pairs read as input by `answerWordCountRequests`.

(g) The number of key-value pairs produced as output by `answerWordCountRequests`.

6. Answer the questions in the collaboration policy on page 1.

4.2 Autolab Implementation Details

Following the first two steps in the algorithm described in the Section 2, you should implement the **Aggregate.java** class, and merge the counts using:

```
cat bigram.txt | sort -k1 | java -Xmx128m Aggregate 1 > bigram_processed.txt
cat unigram.txt | sort -k1 | java -Xmx128m Aggregate 0 > unigram_processed.txt
```

The only argument in the **Aggregate.java** function is the identifier for n-gram: 1 means the file being processed is the bigram counts, while 0 denotes the unigram setup.

According to the third step in the phrase finding algorithm, you should implement and run the **MessageGenerator.java** class, using the following command:

```
cat bigram_processed.txt | java -Xmx128m MessageGenerator > message.txt
```

Next, you will need to write the **MessageUnigramCombiner.java** class that follows the step 4 in the algorithm described in Section 2. The java class must be able to run using the following command:

```
cat message.txt unigram_processed.txt | sort -k1,1 | \
java -Xmx128m MessageUnigramCombiner > message_unigram.txt
```

Finally, in order to generate the final results, you need to write the **PhraseGenerator.java**, which follows the step 5 in the algorithm, and should be able to run using:

```
cat message_unigram.txt bigram_processed.txt | sort -k1,2 | \
java -Xmx128m PhraseGenerator
```

Your **PhraseGenerator** code should write to stdout the **top-20** phrases sorted by by total score. To be more specific, it must print out the phrases and the scores in this format:

```
<phrase>\t<total score>\t<phraseness score>\t<informativeness score>
...
```

where total score is the sum of phraseness score and informativeness score.

You should tar the following items into **hw2.tar** and submit to the homework assignment via Autolab:

- Aggregate.java

- MessageGenerator.java

- MessageUnigramCombiner.java

- PhraseGenerator.java

- and all other auxiliary functions you have written

- report.pdf

Tar the files directly using "tar -cvf hw2.tar *.java report.pdf". Do **NOT** put the above files in a folder and then tar the folder. You do not need to upload the saved temporary files. Please make sure your code is working fine on linux.andrew.cmu.edu machines before you submit.

5   Submission

You must submit your homework through Autolab via the "Homework3: Streaming Phrase Finding" link. In this homework, we provide an additional tool called "Homework3-validation":

- Homework3-validation: You will be notified by Autolab if you can successfully finish your job on the Autolab virtual machines. Note that this is not the place you should debug or develop your **algorithm**. All development should be done on linux.andrew.cmu.edu machines. This is basically a Autolab debug mode. There will be **NO** feedback on your **performance** in this mode. You have unlimited amount of submissions here. To avoid Autolab queues on the submission day, the validation link will be closed 24 hours prior to the official deadline. If you have received a score of 1000, this means that you code has passed the validation.

- Homework3: Streaming Phrase Finding: This is where you should submit your validated final submission. You have a total of **10 possible submissions**. Your performance will be evaluated, and feedback will be provided immediately.

6   Grading

You will be graded based on:

- Memory usage and runtime of the commands that run the aggregation class (10 points), the message generator class (10 points), the message-unigram combiner class (10 points), and the phrase generator class (10 points).

- Precision and recall [1] of top 20 phrases of your output (20 points), and the Spearman correlation[2] of the scores from the top-20 phrases (20 points).

The report will be graded manually (20 points).

---

[1]http://en.wikipedia.org/wiki/Precision_and_recall
[2]http://en.wikipedia.org/wiki/Spearman's_rank_correlation_coefficient

7   Hints/Good to know

Unix sort can handle very large files. This is helpful when you need to collect together the counts for a particular key. To ensure sort behaves properly, you should set the following environment variable:

```
LC_ALL='C'
```

Get it to sort using tabs by setting this flag:

```
-t $'\t'
```

And, when files are large it may be a good idea to tell sort where to store its temporary files:

```
-T /some/dir
```

Note that we have already set up the sort command on Autolab, but you need to check if the above also works in your developmental environment.